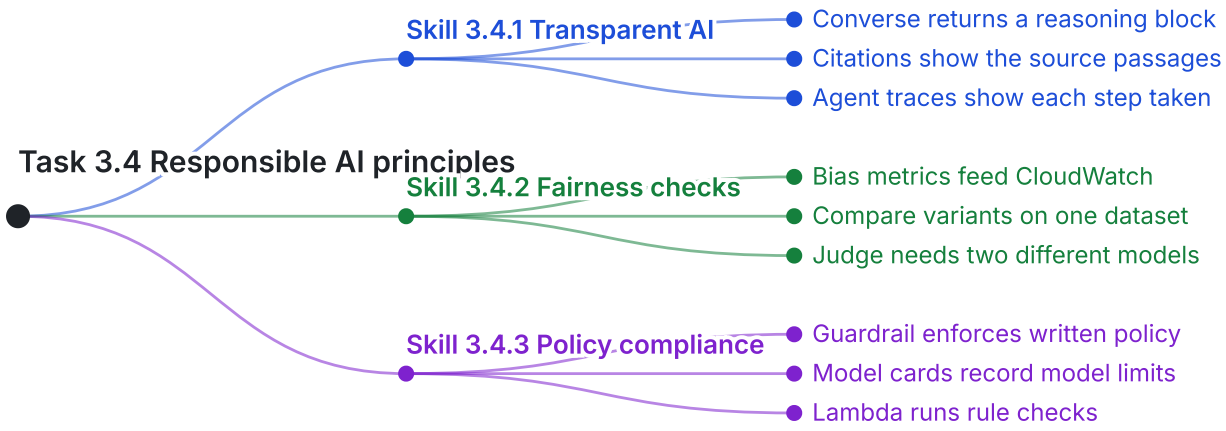
Memory sheet: Task 3.4 Implement responsible AI principles

Domain 3, weight 20 percent. Read this sheet, then answer the flash questions aloud.



Deciding rules

When the stem says	Think of	Not when
Explain why, show reasoning, or show each agent step	Reasoning content block or agent trace	A question about whether the answer is supported, since reasoning does not prove support
Cite, source, or where did this come from	Citations from a knowledge base, using RetrieveAndGenerate	A claim that must be verified, since a citation does not prove the answer is right
Block, mask, refuse, or a policy on prompts	Guardrail attached to the model call	Text that code builds with no model call on that path
Check text that code builds, with no model call	ApplyGuardrail, with the source set to OUTPUT	A model call already on the path, which the guardrail on that call covers
Groups, bias, imbalance, or demographic	Clarify bias metrics published to CloudWatch	One generated paragraph checked for stereotypes, which a judge metric does
A rubric with no answer key on a fixed prompt set	LLM-as-a-judge evaluation job	People who must rate the answers, which needs a human evaluation job
Prompt versions, rollout, or split traffic	Prompt Management variants tested on one dataset	A fairness score, since Prompt Management does not score fairness by itself
Intended use, limitations, or risk rating	SageMaker model card	Runtime enforcement, since a card does not block anything

Etch lines

1. A citation shows where a passage came from, not that it is right.
2. Use ApplyGuardrail when code builds the text and no model call happens.
3. Clarify measures group balance in data and predictions, not single sentences.
4. A judge job scores fixed prompts against a rubric, offline.
5. Prompt Management stores and compares prompt versions but does not score fairness.
6. Lambda rule checks run before the response leaves the system.

Flash questions

- Q: When a reasoning block comes back, what must you send on later Converse requests? A: The signature, with all earlier messages. A changed message makes the request fail.
- Q: Does a citation prove that an answer is correct? A: No. It shows the source passage. Add a grounding check or a judge evaluation to verify the claim.
- Q: Which API lets your code screen a template it fills in, when no model runs on that path? A: ApplyGuardrail, with the source set to OUTPUT.
- Q: Which SageMaker feature reports bias numbers from a live endpoint to CloudWatch? A: Clarify bias drift monitoring, with metrics in a bias-metrics namespace for real-time endpoints.
- Q: Which evaluation fits a stereotype check on one generated paragraph with a fixed test set? A: An LLM-as-a-judge job with the stereotyping metric.
- Q: Which Bedrock feature stores prompt variants and versions that an app calls by ARN? A: Prompt Management. It stores and compares prompts but does not score fairness.
- Q: On InvokeModel, which markup limits the text that a guardrail checks? A: guardContent tags around only the user text, with a new random suffix on each request.
- Q: Which record documents limits and intended use without enforcing anything at runtime? A: A SageMaker model card. Pair it with a guardrail or a deployment gate.

Source: Skill 3.4.1 Develop transparent AI systems in FM outputs, Skill 3.4.2 Apply fairness evaluations to ensure unbiased FM outputs, Skill 3.4.3 Develop policy-compliant AI systems, and The deciding constraint. AWS docs and the study guide sections above, no practice-test text.